



Variational Linear Attention: Constant-Memory Long-Context Processing

(a) KV State Norm ($105\times$ lower at $T=2000$)



(b) Streaming: $\|S_t\|_F$ at T up to 1M (BOUNDED)



(c) GPU Memory vs Sequence Length

