



Variational Linear Attention v3: Empirical Evaluation


(a) KV State Norm ($152\times$ lower at $T=2000$)



(b) MQAR Retrieval ($d=64$, 8 pairs)



(c) Throughput ($d=64$, $B=4$, T4 GPU)



(d) Capacity Degradation ($d=64$)

